



An Introduction to Kubernetes



Cemal Ünal

- Software Engineer @ Havelsan Inc.



cemalunal



cemalunal



cunal@havelsan.com.tr



Container Orchestration

Requirements:

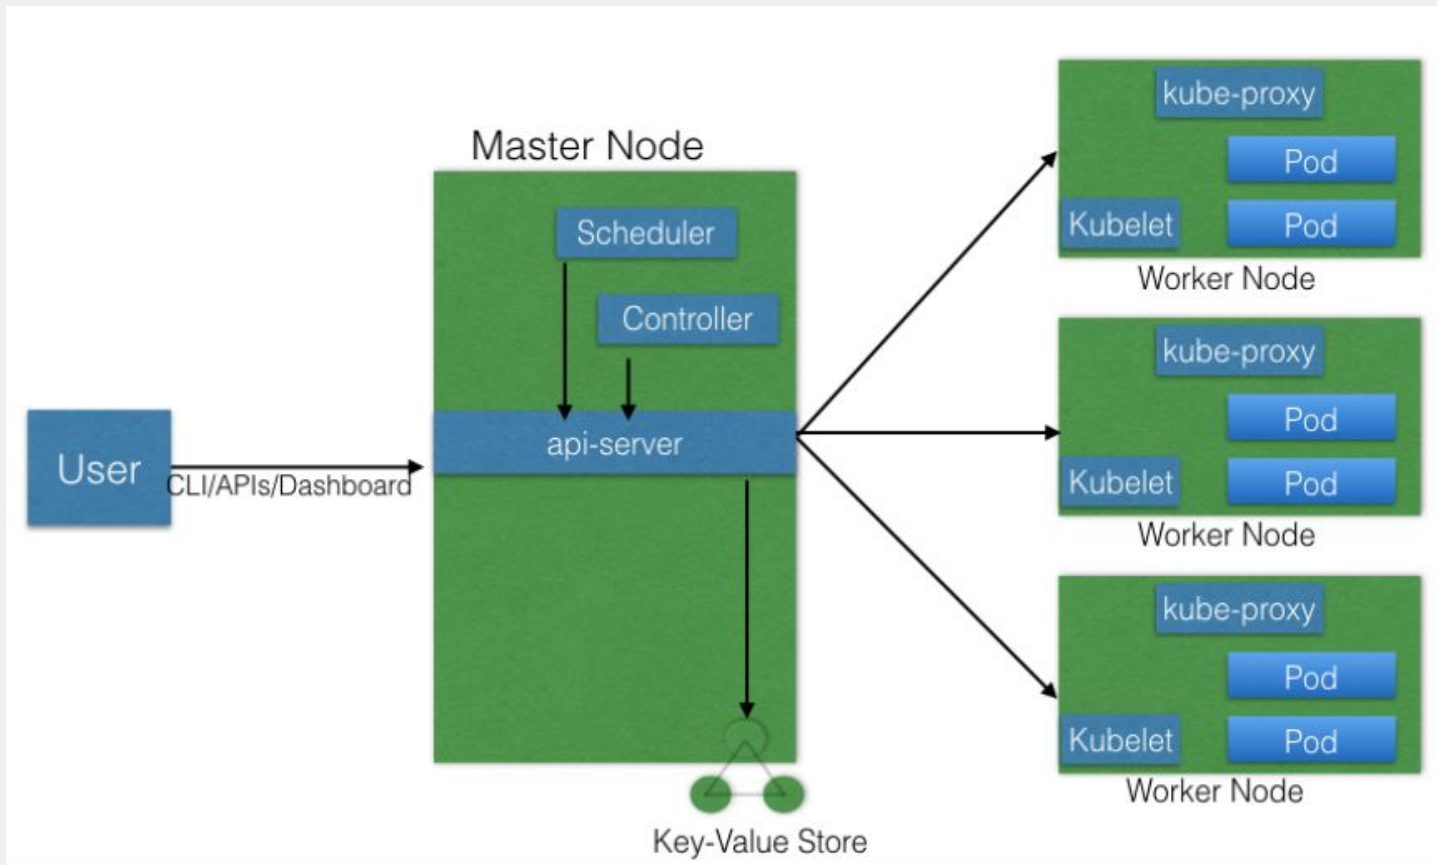
- Fault-tolerance
- On-demand scalability
- Optimal resource usage
- Self healing
- Highly available in case of any failures
- Auto-discovery to automatically discover and communicate with each other
- Seamless updates/rollbacks without any downtime



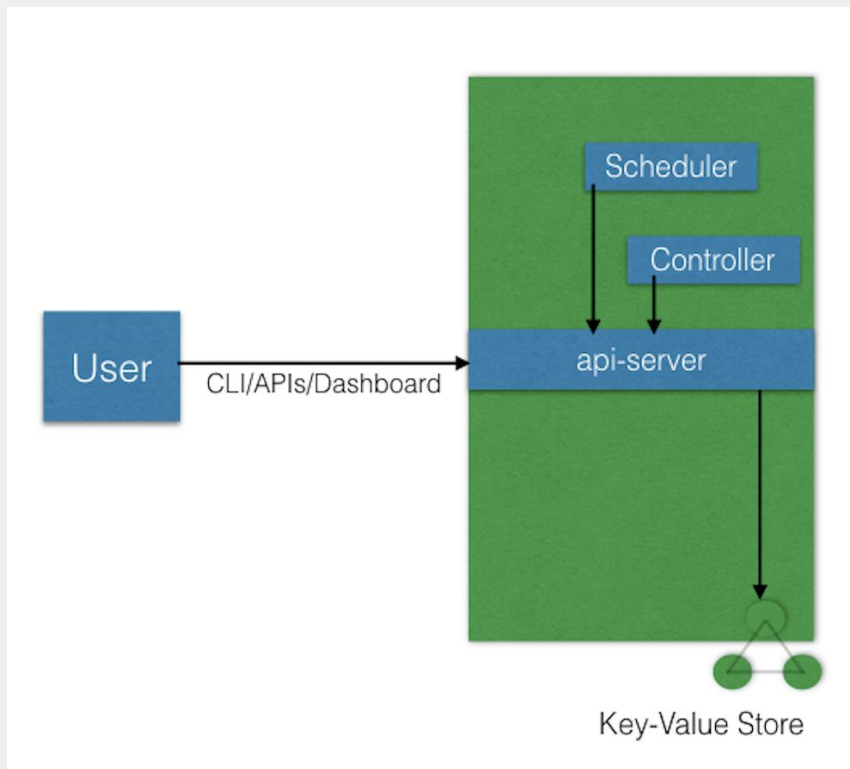
Kubernetes Features

- Self-healing
 - Kills and restarts the containers which do not respond to health checks
- Horizontal scaling
 - Can automatically scale applications based on resource usage like CPU and memory
- Service discovery and Load balancing
 - Groups sets of containers and refers to them via a Domain Name System (DNS). And load-balances requests between these group of containers
- Automated rollouts and rollbacks
 - Can roll out and roll back new versions/configurations of an application, without introducing any downtime

Kubernetes



Kubernetes - Master Node



- Responsible to manage the cluster
- Components
 - API Server
 - Scheduler
 - Component Manager
 - etcd: store the cluster state




Kubernetes Concepts - Pods

- Smallest and simplest Kubernetes object.
- Represents a single instance of the application.
- They do not have the capability to self-heal by themselves. (Use deployment instead)

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.17.0
      ports:
        - containerPort: 80
```

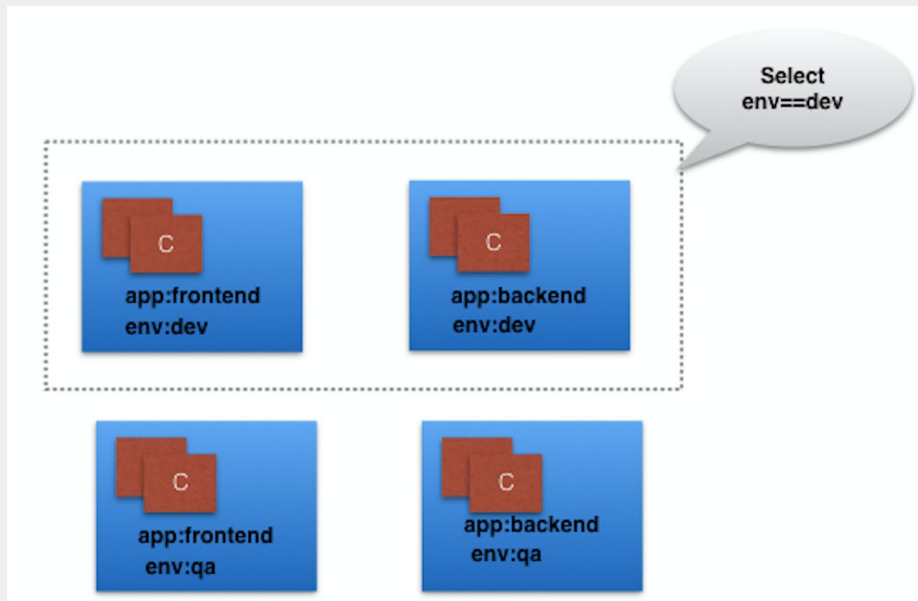
Kubernetes Concepts - Pods

- Are scheduled together on the same host
- Share the same network namespace
- Mount the same external storage (volumes)



Kubernetes Concepts - Labels

- Key-value pairs that can be attached to any Kubernetes object.
- Labels are used to organize and select a subset of objects, based on the requirements in place



Kubernetes Concepts - Deployment

- Makes sure that the current running number of pods always matches the desired state.
- Provide declarative updates to Pods.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  namespace: default
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - image: nginx:1.16.0
        name: nginx
        imagePullPolicy: IfNotPresent
        ports:
        - containerPort: 80
        terminationGracePeriodSeconds: 10
```



Kubernetes Concepts - Namespaces

- Use it to segment the services into manageable chunks.
- Namespaces can make Kubernetes a lot more manageable and gives us increased control, security and flexibility.



```
apiVersion: v1
kind: Namespace
metadata:
  name: elasticsearch-cluster
```

Kubernetes Concepts - Services


- Using Services, Pods can be grouped to provide common access points from the external world.
- Service types:
 - Headless
 - ClusterIP
 - NodePort
 - LoadBalancer

```
kind: Service
apiVersion: v1
metadata:
  name: nginx
  namespace: kube-monitoring
  labels:
    app: nginx
spec:
  type: NodePort
  ports:
    - port: 9093
  selector:
    app: nginx
```

Kubernetes Concepts - Services



Kubernetes Concepts - Services



Kubernetes Concepts - ConfigMap


- Allow us to decouple the configuration details from the container image
- Using ConfigMaps, we can pass configuration details as key-value pairs, which can be later consumed by Pods

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: frontend-cm
  namespace: k8s-demo
data:
  REACT_APP_BACKEND_URI: "http://10.0.0.0"
```

```
spec:
  containers:
    - image: cunal/demo-frontend:v0.0.1
      imagePullPolicy: IfNotPresent
      envFrom:
        - configMapRef:
            name: simple-frontend-cm
```

Horizontal Pod Autoscaler

- Automatically scales the number of pods in a deployment.
- Autoscale is based on observed CPU utilization or some custom metrics.





DEMO



Role Based Access Control (RBAC)

- Mostly useful for
 - Giving access to pods calling Kubernetes API (with Kubernetes Service Account)
 - Giving fine-grained access to people/groups calling Kubernetes API

ClusterRole	A preset of capabilities
Role	ClusterRole, but namespace-scoped
ClusterRoleBinding	Give permission of a ClusterRole to users
RoleBinding	ClusterRoleBinding, but namespace-scoped



Role Based Access Control (RBAC)

Example **Role + Binding** only allows **view of deployments** in the **development** namespace:

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  namespace: development
  name: deployment-viewer
rules:
- apiGroups: ["", "extensions", "apps"]
  resources: ["deployments", "replicasets", "pods"]
  verbs: ["get", "list"]
```

```
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: deployment-viewer-binding
  namespace: development
subjects:
- kind: User
  name: cema1
  apiGroup: ""
roleRef:
  kind: Role
  name: deployment-viewer
  apiGroup: ""
```



Resource Quotas

Limits **memory/cpu/storage** that pods can use, and how many objects of each type (**pods, load balancers, ConfigMaps**) on a **per-namespace** basis

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-resources
  namespace: my-namespace
spec:
  hard:
    requests.cpu: "1"
    requests.memory: 1Gi
    limits.cpu: "2"
    limits.memory: 2Gi
```

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: object-counts
  namespace: my-namespace
spec:
  hard:
    configmaps: "10"
    pods: "4"
    deployments: "20"
    services: "10"
    services.loadbalancers: "2"
```



Assigning Pods to Nodes

```
$ kubectl label nodes node01 disktype=ssd
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    env: test
spec:
  containers:
  - name: nginx
    image: nginx
    imagePullPolicy: IfNotPresent
  nodeSelector:
    disktype: ssd
```



THANKS!