



# **An Introduction to Prometheus**



# Cemal Ünal

- Software Engineer @ Havelsan Inc.



cemalunal



cemalunal



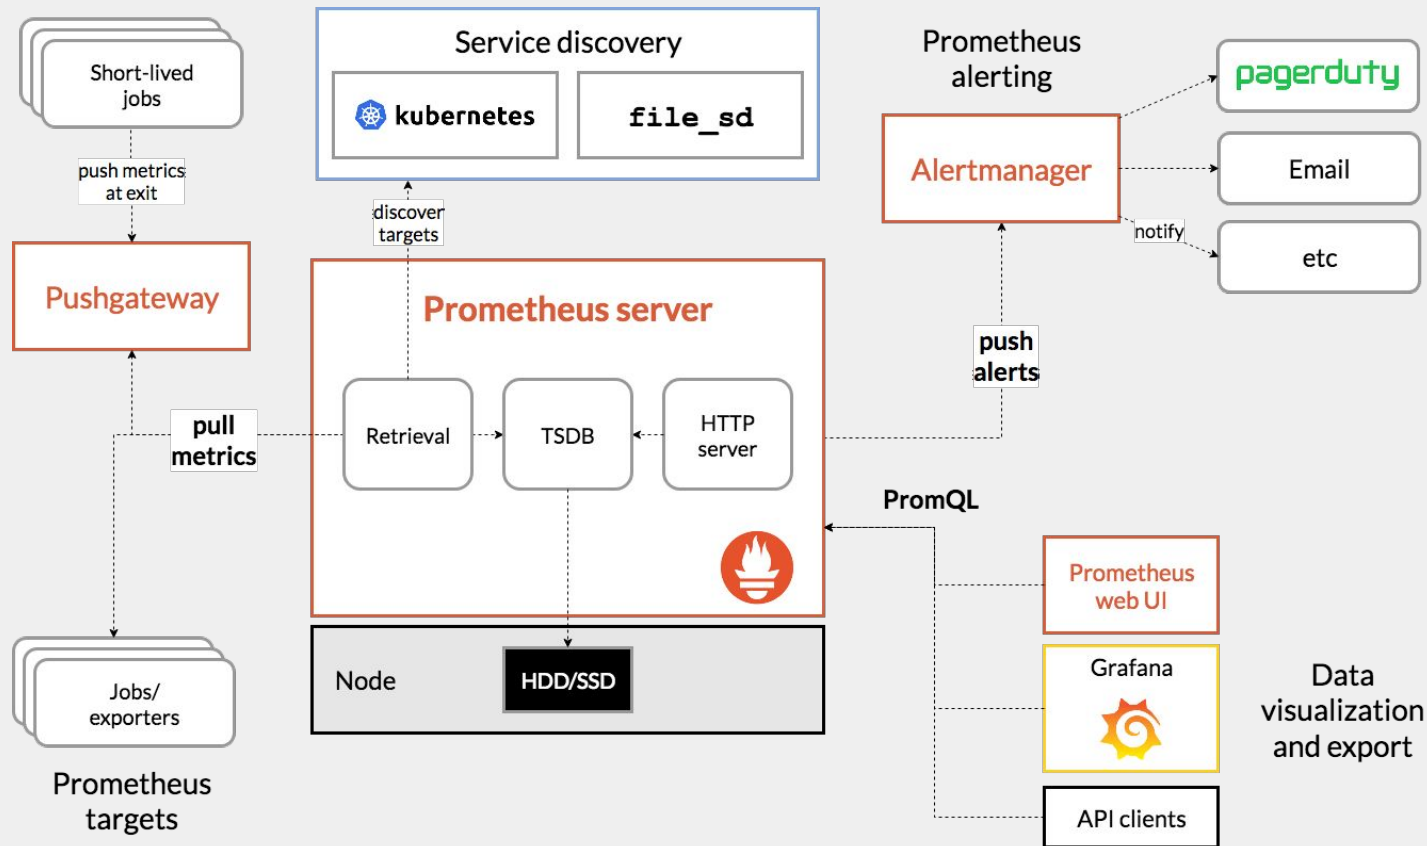
cunal@havelsan.com.tr

# Prometheus

- Open-source systems monitoring and alerting toolkit
- Records real-time metrics in a time series database built using a HTTP pull model
- Metric collection, processing and storage
- Powerful query language
- Metric Visualization
- Real time alerting



# Ecosystem and Architecture





# Components

The ecosystem consists of multiple components. Some of them are:

- The main **Prometheus server** which scrapes and stores time series data
- **Special-purpose exporters** for services like HAProxy, StatsD, Graphite, etc.
- **AlertManager** to handle alerts
- **Push Gateway** supporting short living jobs



# Data Model

- Fundamentally stores all data as time series (measurement/metric with a timestamp).
- Each time series is uniquely identified by a metric name and a set of key-value pairs (labels).
- `<metric name>{<label name>=<label value>, ...}`



# Metric Types

There are four different types of metrics:

- **Counter:** A cumulative metric that only ever increases. (E.g. requests served, tasks completed, errors occurred)
- **Gauge:** A metric that can arbitrarily go up or down. (E.g. temperature, current memory usage)
- **Histogram:** Binned measurement of a continuous variable. (E.g. latency, request duration)
- **Summary:** Similar to a histogram, except the bins are converted into an aggregate (e.g. 99% percentile) immediately.



# Metric Example

```
http_requests_total{service="webserver", region="EU"} 136
```

- Metric name is **http\_requests\_total**
- Two labels, **service** and **region** with values "**webserver**" and "**EU**"
- Metric value is **136**

This webserver service has received 136 http requests since its started.





# PromQL Examples

## **Filter:**

```
http_requests_total{service="webserver"}
```

*Total HTTP requests that webserver receives*

## **Aggregate:**

```
sum(http_requests_total)
```

*total HTTP requests*

## **Average rate of increase:**

```
rate(http_requests_total[5m])
```

*the per-second rate of HTTP requests as measured over the last 5 minutes*



# Alerting

Prometheus allows defining rules based on PromQL in the following format:

- `alert: LowRequests`  
    `expr: rate(http_requests_total{service="webserver"}[1m]) < 0.1`  
    `for: 5m`

Prometheus will execute this regularly.

AlertManager will send a notification via email, webhook, etc.



# Visualization - Built-in

- Allows us to enter any expression and see its result either in a table or graphed over time.
- Useful for ad-hoc queries and debugging purposes
- Available at `/graph` endpoint

☐ Enable query history[Try experimental React UI](#)

```
rate(node_cpu_seconds_total{mode="system"}[1m])
```


Execute

- insert metric at cursor - ▾

Load time: 36ms  
Resolution: 14s  
Total time series: 12

Graph Console

- 1h + ◀ Until ▶ Res. (s) ☐ stacked

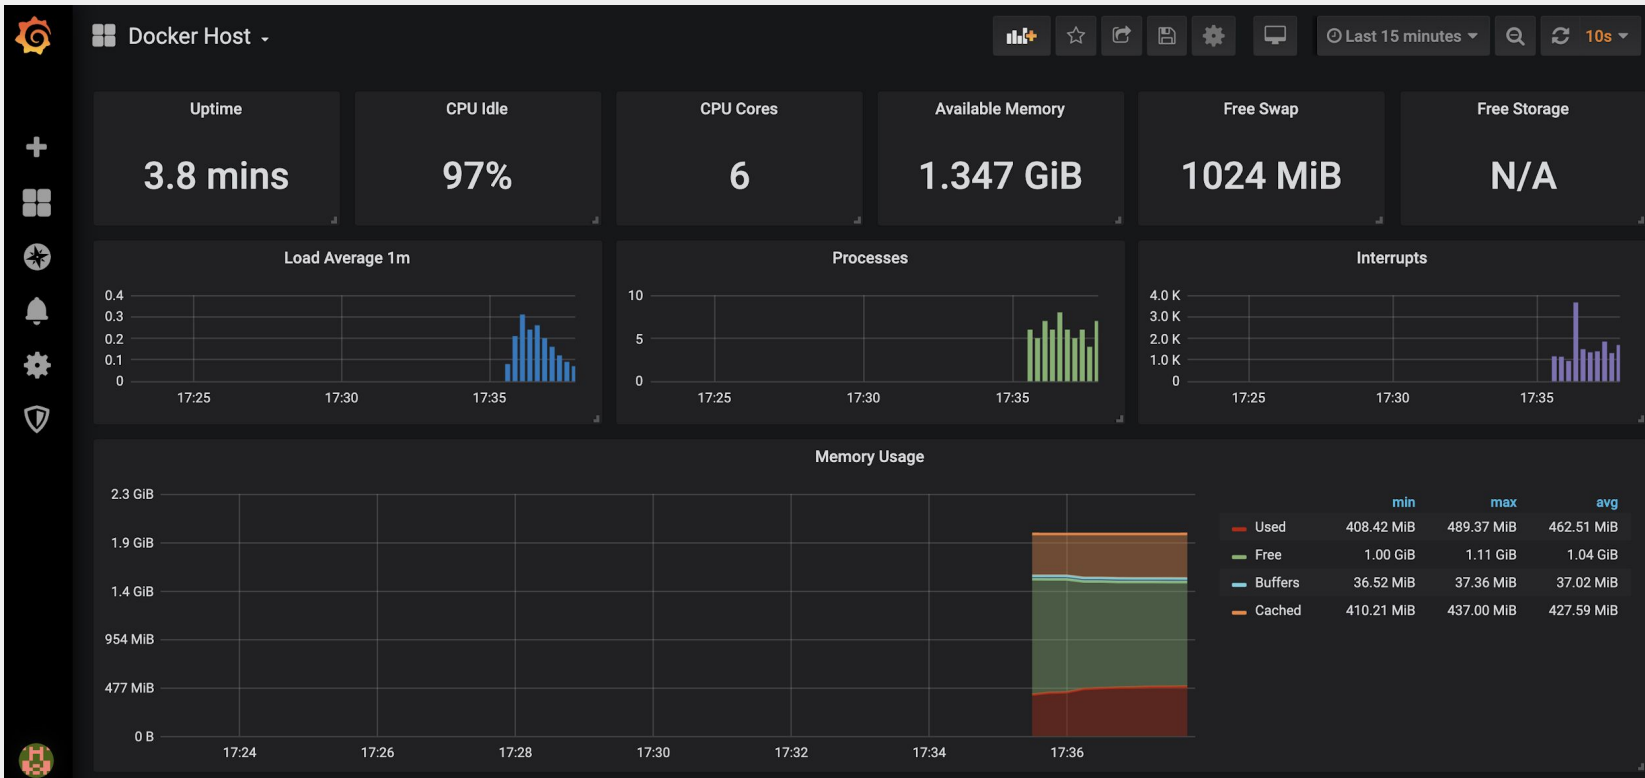


✓ {cpu="9",instance="localhost:9100",job="node-exporter",mode="system"}  
✓ {cpu="8",instance="localhost:9100",job="node-exporter",mode="system"}



# Visualization - Grafana

- Open source analytics and monitoring solution
- Supports querying Prometheus metrics





# DEMO



**THANKS!**